

Homework 2 Solutions

October 11, 2018

Part 1

i. Since *NYChousing* is a .csv file I use *read.csv()* to import the data into R.

```
setwd("~/Desktop/Data")
housing <- read.csv("NYChousing.csv", as.is = TRUE)
```

ii. The function *dim()* provides the dimension of its input object.

```
orig_dim <- dim(housing)
orig_dim
```

```
## [1] 2506 22
```

iii.

```
apply(is.na(housing), 2, sum)
```

```
##                UID                PropertyName
##                0                        0
##                Lon                    Lat
##                15                    15
##                AgencyID                Name
##                0                        0
##                Value                    Address
##                52                        0
##                Violations2010            REACNumber
##                0                        1873
##                Borough                  CD
##                0                        0
##                CityCouncilDistrict        CensusTract
##                10                        0
##                BuildingCount              UnitCount
##                0                        0
##                YearBuilt                  Owner
##                0                        0
##                Rental.Coop              OwnerProfitStatus
##                0                        0
##                AffordabilityRestrictions StartAffordabilityRestrictions
##                0                        5
```

The command *is.na(housing)* creates a matrix of the same dimensions as *housing* with each element being TRUE or FALSE depending on whether or not the corresponding element in *housing* is an NA value. Then the full call *apply(is.na(housing), 2, sum)* counts the number of NA values each column of *housing*.

iv.

```
housing <- housing[!is.na(housing$Value), ]
```

The call *is.na(housing\$Value)* returns a logical vector with TRUE where *housing\$Value* is NA, therefore I filter using *!is.na(housing\$Value)* to get only the rows where *Value* is not NA. I reassign my *housing* dataframe, to be the filtered dataframe.

v.

```
new_dim <- dim(housing)
orig_dim[1] - new_dim[1]
```

```
## [1] 52
```

I removed 52 rows of my dataframe which is what I expect, since my output in (iii) told me that I have 52 missing values in **Value**.

v.

```
housing$logValue <- log(housing$Value)
summary(housing$logValue)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      8.41  12.49   13.75   13.68   14.80   20.47
```

vi.

```
housing$logUnits <- log(housing$UnitCount)
```


vii.

```
housing$after1950 <- housing$YearBuilt >= 1950
```

Part 2: EDA

i.


```
plot(housing$logUnits, housing$logValue, xlab = "log(Units)", ylab = "Log(Value)")
```



I plot a scatterplot with the **plot()** command and add argument **xlab =** and **ylab =** for the labels.

ii.

```
plot(housing$logUnits, housing$logValue, col = factor(housing$after1950), xlab = "log(Units)", ylab = "Log(Value)",
legend("bottomright", legend = levels(factor(housing$after1950)), fill = unique(factor(housing$after1950)))
```



There appears to be a pretty strong linear relationship between **logValue** and **logUnits**. When colored according to the **after1950** variable, it is clear that newer buildings (those built after 1950) tend to be more expensive and have more units than older buildings.

iii.

```
cor(housing$logValue, housing$logUnits)
```

```
## [1] 0.8727348
```

```
cor(housing$logValue[housing$Borough == "Manhattan"], housing$logUnits[housing$Borough == "Manhattan"])
```

```
## [1] 0.8830348
```

```
cor(housing$logValue[housing$Borough == "Brooklyn"], housing$logUnits[housing$Borough == "Brooklyn"])
```

```
## [1] 0.9102601
```

```
cor(housing$logValue[housing$after1950], housing$logUnits[housing$after1950])
```


```
## [1] 0.721735
```

```
cor(housing$logValue[!housing$after1950], housing$logUnits[!housing$after1950])
```

```
## [1] 0.8643297
```

iv.

```
plot(housing$logUnits[housing$Borough == "Manhattan"], housing$logValue[housing$Borough == "Manhattan"])
points(housing$logUnits[housing$Borough == "Brooklyn"], housing$logValue[housing$Borough == "Brooklyn"])
legend("bottomright", legend = c("Manhattan", "Brooklyn"), fill = c("black", "red"))
```



v.

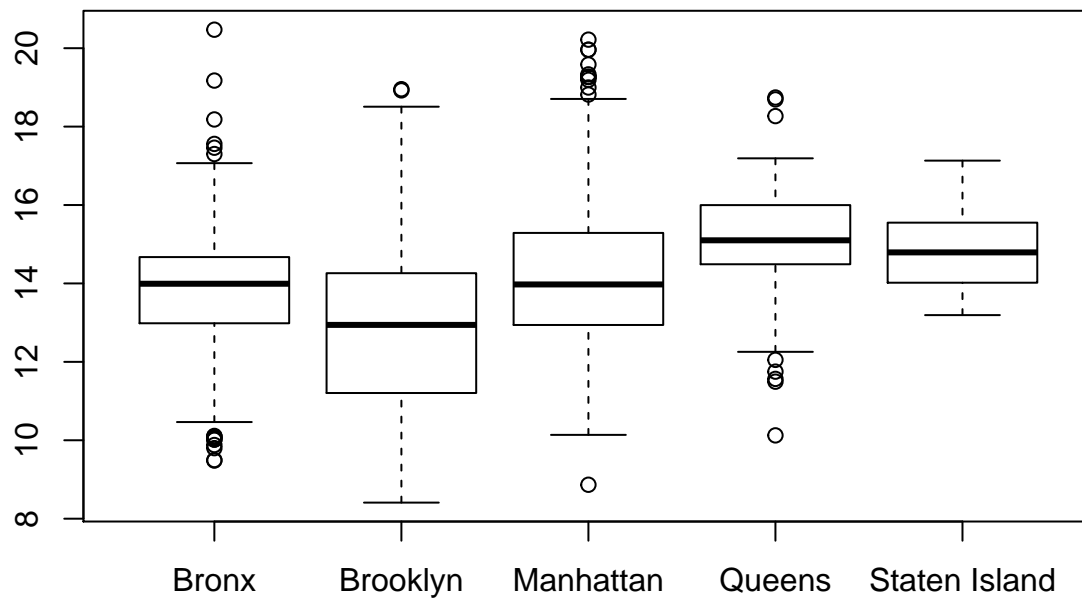
```
median(housing$Value[housing$Borough == "Manhattan"])
```

```
## [1] 1172362
```

The code calculates the median property value for all properties in Manhattan.

vi.

```
boxplot(housing$logValue ~ housing$Borough)
```



vii.

```
tapply(housing$Value, housing$Borough, median)
```

##	Bronx	Brooklyn	Manhattan	Queens	Staten Island
##	1192950	417610	1172362	3611700	2654100

We use **tdapply()** which splits the property value into groups based on **Borough** and then calculated the median within each group.